

Project task 1:

Data import and preparation:

```
## importing libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

## importing data

train= pd.read_csv('train.csv')
test= pd.read_csv('test.csv')

df= pd.concat([test,train])

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 39030 entries, 0 to 27320
Data columns (total 80 columns):
#   Column                                Non-Null Count  Dtype
---  -
0    UID                                39030 non-null  int64
1    BLOCKID                           0 non-null     float64
2    SUMLEVEL                          39030 non-null  int64
3    COUNTYID                          39030 non-null  int64
4    STATEID                           39030 non-null  int64
5    state                             39030 non-null  object
6    state_ab                           39030 non-null  object
7    city                             39030 non-null  object
8    place                             39030 non-null  object
9    type                             39030 non-null  object
10   primary                           39030 non-null  object
11   zip_code                          39030 non-null  int64
12   area_code                         39030 non-null  int64
13   lat                               39030 non-null  float64
14   lng                               39030 non-null  float64
15   ALand                             39030 non-null  float64
16   AWater                            39030 non-null  int64
17   pop                               39030 non-null  int64
18   male_pop                          39030 non-null  int64
19   female_pop                        39030 non-null  int64
20   rent_mean                         38568 non-null  float64
21   rent_median                       38568 non-null  float64
22   rent_stdev                        38568 non-null  float64
23   rent_sample_weight               38568 non-null  float64
24   rent_samples                     38568 non-null  float64
25   rent_gt_10                       38567 non-null  float64
26   rent_gt_15                       38567 non-null  float64
27   rent_gt_20                       38567 non-null  float64
28   rent_gt_25                       38567 non-null  float64
29   rent_gt_30                       38567 non-null  float64
30   rent_gt_35                       38567 non-null  float64
31   rent_gt_40                       38567 non-null  float64
32   rent_gt_50                       38567 non-null  float64
33   universe_samples                 39030 non-null  int64
34   used_samples                     39030 non-null  int64
35   hi_mean                           38640 non-null  float64
36   hi_median                        38640 non-null  float64
37   hi_stdev                         38640 non-null  float64
38   hi_sample_weight                 38640 non-null  float64
39   hi_samples                       38640 non-null  float64
40   family_mean                      38596 non-null  float64
41   family_median                    38596 non-null  float64
42   family_stdev                     38596 non-null  float64
43   family_sample_weight             38596 non-null  float64
44   family_samples                   38596 non-null  float64
45   hc_mortgage_mean                 38189 non-null  float64
46   hc_mortgage_median               38189 non-null  float64
47   hc_mortgage_stdev                38189 non-null  float64
48   hc_mortgage_sample_weight        38189 non-null  float64
49   hc_mortgage_samples              38189 non-null  float64
50   hc_mean                          38140 non-null  float64
51   hc_median                        38140 non-null  float64
52   hc_stdev                         38140 non-null  float64
```

	UID	BLOCKID	SUMLEVEL	COUNTYID	STATEID	state	state_ab	city	place	type	...	female_age_mean	female_
0	255504	NaN	140	163	26	Michigan	MI	Detroit	Dearborn Heights City	CDP	...	34.78682	
1	252676	NaN	140	1	23	Maine	ME	Auburn	Auburn City	City	...	44.23451	
2	276314	NaN	140	15	42	Pennsylvania	PA	Pine City	Millerton	Borough	...	41.62426	
3	248614	NaN	140	231	21	Kentucky	KY	Monticello	Monticello City	City	...	44.81200	
4	286865	NaN	140	355	48	Texas	TX	Corpus Christi	Edroy	Town	...	40.66618	
5 rows × 80 columns													

```
## checking for null values (in percentage %)

df.isnull().sum()*100/ len(df)
```

UID	0.000000
BLOCKID	100.000000
SUMLEVEL	0.000000
COUNTYID	0.000000
STATEID	0.000000
...	
pct_own	0.999231
married	0.704586
married_snp	0.704586

```
separated      0.704586
divorced       0.704586
Length: 80, dtype: float64

## number of null values in numbers

df.isnull().sum()

UID            0
BLOCKID       39030
SUMLEVEL      0
COUNTYID     0
STATEID       0
...
pct_own        390
married        275
married_snp    275
separated      275
divorced       275
Length: 80, dtype: int64

## removing null column

df.drop('BLOCKID', axis=1,inplace=True)
df.dropna(inplace=True)

## filling null values with mean of their respective column

for i in df.columns:
    if df[i].isnull().sum()!=0:
        df.fillna(df[i].mean(), inplace=True)
```

```
## checking for null values

df.isna().sum()

UID            0
SUMLEVEL      0
COUNTYID     0
STATEID       0
state         0
..
pct_own        0
married        0
married_snp    0
separated      0
divorced       0
Length: 79, dtype: int64
```

2,500 locations where the percentage of households with a second mortgage is the highest and percent ownership is above 10 percent.

```
top=df.nlargest(2500,'second_mortgage')


top= top.loc[:,['state','city','zip_code','area_code','lat','lng','second_mortgage']]
top.head()
```

	state	city	zip_code	area_code	lat	lng	second_mortgage
14014	New Jersey	Passaic	7055	973	40.867944	-74.114633	0.60870
6238	New York	Bronx	10452	718	40.842166	-73.926952	0.58824
3285	Virginia	Farmville	23901	434	37.297357	-78.396452	0.50000
21706	Arizona	Scottsdale	85257	480	33.458658	-111.955104	0.43750
11980	Massachusetts	Worcester	1610	508	42.254262	-71.800347	0.43363

```
import plotly.express as px
```

Scatter geo plot for top 2500 locations with highest second mortgage

```
px.scatter_geo(top, lat='lat',lon='lng',color= 'second_mortgage',hover_name= 'state',projection='mollweide' )
```



The geoplots shows scattered dots for top 2500 locations in given dataset having highest second mortgage

▼ Bad debt calculations

```
bad_debt= (df.second_mortgage + df.home_equity) - df.home_equity_second_mortgage
```


```
bad_debt

0      0.07651
1      0.14375
2      0.06744
3      0.01741
4      0.03440
...
27316   0.00000
27317   0.20908
27318   0.07857
27319   0.14305
27320   0.18362
Length: 37940, dtype: float64
```

```
overall_debt= df['second_mortgage'] + df['home_equity'] + bad_debt
```

```
d1= [overall_debt.sum(), bad_debt.sum()]
l1= ['overall debt', 'bad debt']
```

```
plt.pie(d1, labels= l1, explode= [0,0.05], autopct='%1.1f%%')
plt.show()
```



```
## adding bad debt column to our original dataset
```

```
df['bad_debt']= bad_debt
```

```
df.bad_debt.median()
```

```
0.09961
```

▼ Box and whisker plot and for 2nd mortgage, home equity, good debt, and bad debt for cities with highest population

```
df.nlargest(4, 'pop')['city']
```

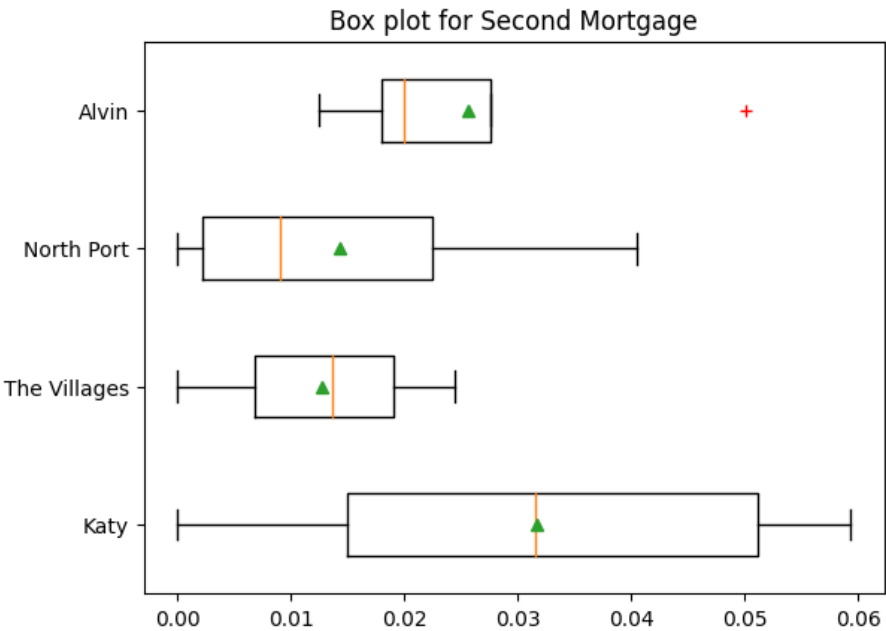
```
169      Katy
23985  The Villages
23565  North Port
15940    Alvin
Name: city, dtype: object
```

The yellow line in the interquartile range is the median

The green triangle in the interquartile range is the mean

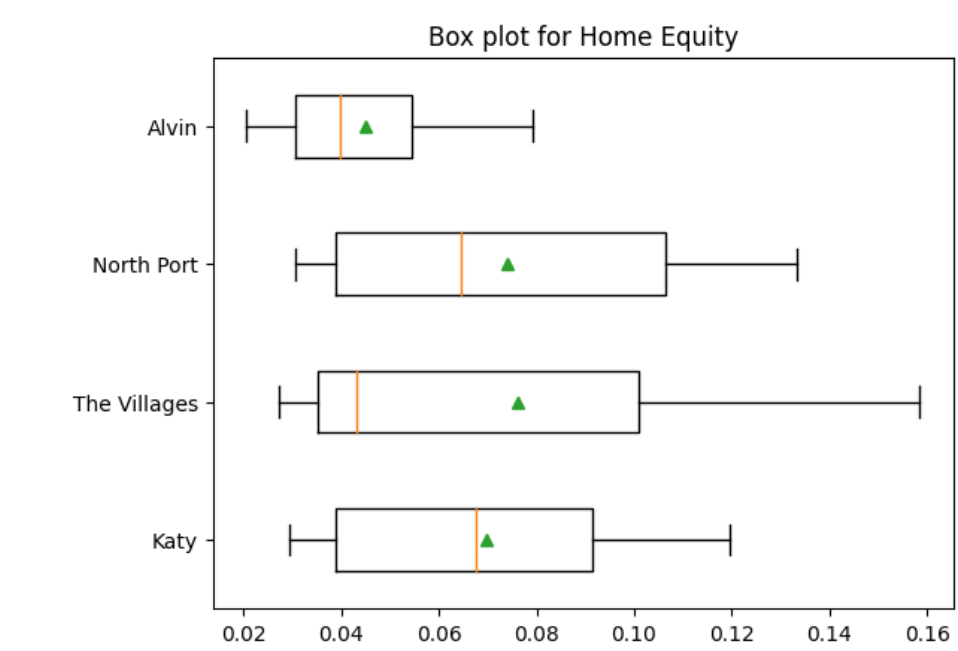
The '+' sign shows the presence of outliers in dataset

```
det= df.loc[df['city']== 'Katy','second_mortgage'].values
las= df.loc[df['city']== 'The Villages','second_mortgage'].values
pine= df.loc[df['city']== 'North Port','second_mortgage'].values
aub= df.loc[df['city']== 'Alvin','second_mortgage'].values
plt.boxplot([det,las,pine,aub],labels=['Katy','The Villages','North Port','Alvin'],showmeans=True, sym= 'r+',vert=False);
plt.title('Box plot for Second Mortgage')
plt.show()
```



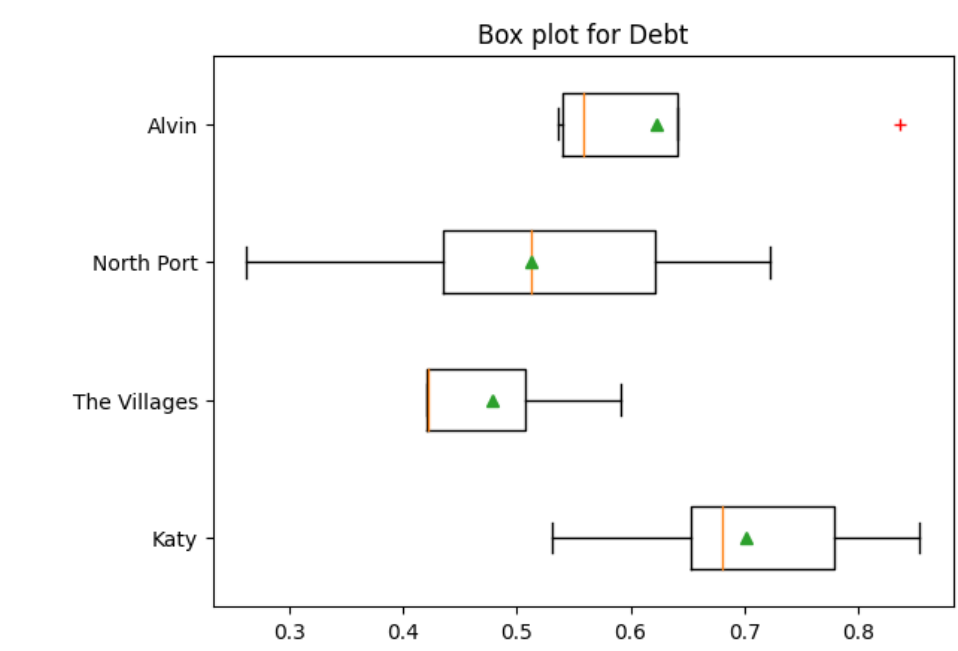
- From the above box plot we can observe the median value of the **second mortgage** data is much closer to the first quartile than the third quartile, which means the **distribution of Alvin, North Port and Katy is right-skewed**.
- From the above box plot we can observe the median value of the **second mortgage** data is much closer to the third quartile than the first quartile which means the **distribution of The Villages is left=skewed**.
- The data contains **outliers above upper quartile range** in the second mortgage data in **Alvin**, which means some people living in Alvin have much more second mortgage as compared to other population in Alvin.

```
was= df.loc[df['city']== 'Katy','home_equity'].values
gf= df.loc[df['city']== 'The Villages','home_equity'].values
bell= df.loc[df['city']== 'North Port','home_equity'].values
clark= df.loc[df['city']== 'Alvin','home_equity'].values
plt.boxplot([was,gf,bell,clark],labels=['Katy','The Villages','North Port','Alvin'],showmeans=True, sym= 'r+',vert=False);
plt.title('Box plot for Home Equity')
plt.show()
```




- From the above box plot we can see the median of the **home equity** data is much closer to the first quartile than the third quartile, which means the **distribution for Alvin, North Port, The Villages is right-skewed**.
- From the above box plot we can observe the median value of the **second mortgage** data is much closer to the third quartile than the first quartile which means the **distribution of Katy is left=skewed**.

```
seat= df.loc[df['city']== 'Katy','debt'].values
colu= df.loc[df['city']== 'The Villages','debt'].values
liver= df.loc[df['city']== 'North Port','debt'].values
frank = df.loc[df['city']== 'Alvin','debt'].values
plt.boxplot([seat,colu,liver,frank],labels=['Katy','The Villages','North Port','Alvin'],showmeans=True, sym= 'r+',vert=False);
plt.title('Box plot for Debt')
plt.show()
```

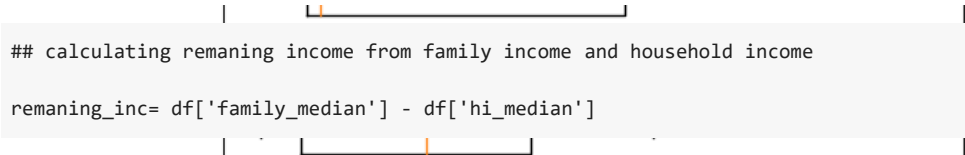


- From the above box plot we can see the median value of **debt** is much closer to the first quartile than the third quartile, which means the **distribution for Alvin, North Port, The Villages and Katy is right skewed**.
- The dataset contains **outliers in debt data above the upper quartile range** for Alvin, which means there are people with very high debt as compared to the median/mean debt of people living in Alvin.

```
ofx= df.loc[df['city']== 'Katy','bad_debt'].values
brook= df.loc[df['city']== 'The Villages','bad_debt'].values
chig= df.loc[df['city']== 'North Port','bad_debt'].values
ang= df.loc[df['city']== 'Alvin','bad_debt'].values
plt.boxplot([ofx,brook,chig,ang],labels=['Katy','The Villages','North Port','Alvin'],showmeans=True, sym= 'r+',vert=False);
plt.title('Box plot for Bad Debt')
plt.show()
```



- From the above box plot we can see the median value of bad debt is much closer to the third quartile than the first quartile, which means the **distribution for Katy is left skewed**.
- From the above box plot we can see the median of the **home equity** data is much closer to the first quartile than the third quartile, which means the **distribution for Alvin, North Port and The Villages is right-skewed**



```
## calculating remaning income from family income and household income

remaning_inc= df['family_median'] - df['hi_median']

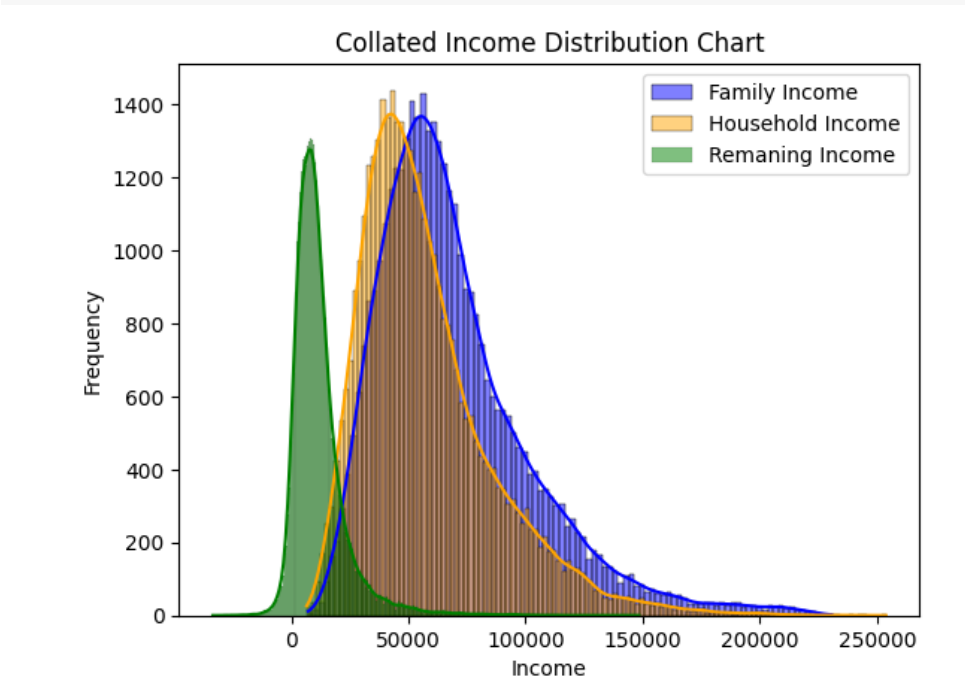
## comparing family income and househould income

df['family_median'].sum() > df['hi_median'].sum()

True
```

```
sns.histplot(df['family_median'],kde=True, color='blue', label='Family Income',legend=True)
sns.histplot(df['hi_median'],kde=True, color='orange', label='Household Income',legend=True)
sns.histplot(remaning_inc,kde=True, color='green', label='Remaning Income',legend=True)
## kde parameter is kernel density estimate to smooth the distribution.
```


```
plt.title('Collated Income Distribution Chart')
plt.xlabel('Income')
plt.ylabel('Frequency')
plt.legend()
plt.show()
```



- We can observe from above collated distribution chart the overall family income and household income lies between 50,000 to 1,00,000 for the given population
- We can observe that the family income is slightly higer than household income for the given population
- The distribution is similar to normal distribution

```
# plt.hist(df['family_median'], bins=75, alpha=0.6, label='Family Income')
# plt.hist(df['hi_median'], bins=75, alpha=0.6, label='Household Income')
# plt.hist(remaning_inc, bins=75, alpha=0.6, label='Remaining Income')
### The alpha parameter is used for transparency to make overlapping histograms more visible.
```

```
# plt.title('Collated Income Distribution')
# plt.xlabel('Income')
# plt.ylabel('Frequency')
# plt.legend()
```



▼ Project task 2:

Exploratory Data Analysis (EDA):

Double-click (or enter) to edit

```
## calculating density of area

density= df['pop']/df['ALand']
```

```
## adding density to original dataset

df['density']= density*100000
df['density']
```

0	126.029034
1	25.685467
2	1.523347
3	0.499905
4	45.157830
...	
27316	264.981421
27317	81.834237
27318	0.213790
27319	61.879771

```
27320      47.791852
Name: density, Length: 37940, dtype: float64
```

```
density.nlargest(3)
```

```
2677      0.175563
21050     0.072283
10251     0.071976
dtype: float64
```

```
df.drop(2677, inplace=True)
```

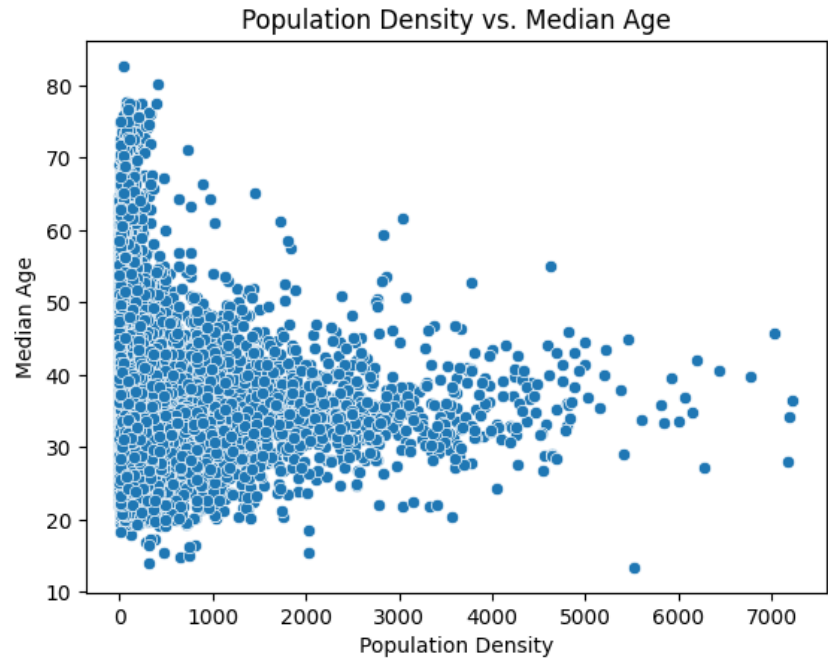
```
## calculating overall median age of male and female population
```

```
median_age_c= ((df['male_age_median']* df['male_pop']) + (df['female_age_median'] * df['female_pop'])) / (df['male_pop']+df['female_pop'])
```

```
df['median_age']= median_age_c
```

```
## Scatter plot of Population density vs. Age
```

```
sns.scatterplot(x='density', y='median_age', data=df)
plt.title('Population Density vs. Median Age')
plt.xlabel('Population Density')
plt.ylabel('Median Age')
plt.show()
```



```
## summary stats into population density and age
```

```
print("Summary Stats: \n")
df[['pop', 'ALand', 'density', 'median_age']].describe()
```

```
Summary Stats:
```

	pop	ALand	density	median_age
count	37940.000000	3.794000e+04	37940.000000	37940.000000
mean	4385.977570	1.251229e+08	200.150973	39.321703
std	2084.057931	1.158857e+09	433.836605	7.397916
min	38.000000	8.299000e+03	0.001172	13.378362
25%	2956.000000	1.824246e+06	12.790635	34.284358
50%	4106.000000	4.951182e+06	86.462514	39.272186
75%	5470.250000	3.453241e+07	206.700889	43.875669
max	53812.000000	1.039510e+11	17556.332088	82.664697

```
## creating bins for age
```


```
pop_bins= pd.cut(df['pop'],5)
pop_bins
```

```
0      (-15.774, 10792.8]
1      (-15.774, 10792.8]
2      (-15.774, 10792.8]
3      (-15.774, 10792.8]
4      (-15.774, 10792.8]
...
27316  (-15.774, 10792.8]
27317  (-15.774, 10792.8]
27318  (-15.774, 10792.8]
27319  (10792.8, 21547.6]
27320  (-15.774, 10792.8]
Name: pop, Length: 37940, dtype: category
Categories (5, interval[float64, right]): [(-15.774, 10792.8] < (10792.8, 21547.6] <
(21547.6, 32302.4] < (32302.4, 43057.2] <
(43057.2, 53812.0]]
```

```
male= df['male_age_median'].median()
female= df['female_age_median'].median()
overall= median_age_c.median()
data1= [male,female ,overall]
label= ['male','female','overall']
```

visulaizing overall, female and male median age for given population

```
plt.bar(label,data1, color=['blue','pink','green'])
plt.title('Median age comparison')
plt.xlabel('Gender')
plt.ylabel('Median Age')
plt.ylim(0, max(data1) + 7)
plt.text(0, male+1 , f'Male Median Age: {male}', ha='center', va='bottom', color='black')
plt.text(1, female+2, f'Female Median Age: {female}', ha='center', va='bottom', color='black')
plt.text(2, overall+1, f'Overall Median Age: {overall:.2f}', ha='center', va='bottom', color='black')
plt.show()
```



- The overall population median age is 39.27 years
- The male population median age is 38 years
- The female population median age is 38 years

df['state'].unique()

```
array(['Michigan', 'Maine', 'Pennsylvania', 'Kentucky', 'Texas',
      'Florida', 'Georgia', 'New York', 'California', 'Washington',
      'Illinois', 'Massachusetts', 'Maryland', 'Virginia', 'Nevada',
      'Tennessee', 'District of Columbia', 'Arkansas', 'Alabama',
      'Wisconsin', 'New Mexico', 'Mississippi', 'Ohio', 'Indiana',
      'Montana', 'Oregon', 'New Jersey', 'North Carolina', 'Louisiana',
      'South Carolina', 'Utah', 'Arizona', 'Rhode Island', 'Puerto Rico',
      'Oklahoma', 'Missouri', 'Minnesota', 'Nebraska', 'Colorado',
      'Iowa', 'Connecticut', 'Delaware', 'Kansas', 'West Virginia',
      'Vermont', 'Alaska', 'South Dakota', 'Idaho', 'New Hampshire',
      'Hawaii', 'Wyoming', 'North Dakota'], dtype=object)
```

```
overall_inc= df['family_median']
overall_rent= df['rent_median']
overall_rent_percent= (overall_rent / overall_inc *100 ).median()
```

```
michigan_inc = df.loc[df['state']=='Michigan','family_median']
michigan_rent = df.loc[df['state']=='Michigan','rent_median']
michigan_rent_percent= (michigan_rent/michigan_inc * 100).median()
```

```
texas_inc = df.loc[df['state']=='Texas','family_median']
texas_rent = df.loc[df['state']=='Texas','rent_median']
texas_rent_percent= (texas_rent/texas_inc * 100).median()
```


```
nyc_inc = df.loc[df['state']=='New York','family_median']
nyc_rent = df.loc[df['state']=='New York','rent_median']
nyc_rent_percent= (nyc_rent/nyc_inc * 100).median()
```

```
indiana_inc = df.loc[df['state']=='Indiana','family_median']
indiana_rent = df.loc[df['state']=='Indiana','rent_median']
indiana_rent_percent= (indiana_rent/indiana_inc * 100).median()
```

```
wyoming_inc = df.loc[df['state']=='Wyoming','family_median']
wyoming_rent = df.loc[df['state']=='Wyoming','rent_median']
wyoming_rent_percent= (wyoming_rent/wyoming_inc * 100).median()
```

```
data2= [overall_rent_percent, michigan_rent_percent, nyc_rent_percent,indiana_rent_percent,wyoming_rent_percent]
labels= ['Overall','Michigan','NewYork','Indiana','Wyoming']
```

```
plt.bar(labels, data2,color=['red','blue','green','yellow','magenta'])
plt.title('Rent Percentage Distribution')
plt.xlabel('States')
plt.ylabel('Rent Percentage')
plt.ylim(0, max(data2)+0.3)
plt.text(0, overall_rent_percent+0.1 , overall_rent_percent.round(3), ha='center', va='bottom', color='black')
plt.text(1, michigan_rent_percent+0.1, michigan_rent_percent.round(3), ha='center', va='bottom', color='black')
plt.text(2, nyc_rent_percent+0.1, nyc_rent_percent.round(3), ha='center', va='bottom', color='black')
plt.text(3, indiana_rent_percent+0.1, indiana_rent_percent.round(3), ha='center', va='bottom', color='black')
plt.text(4, wyoming_rent_percent+0.1, wyoming_rent_percent.round(3), ha='center', va='bottom', color='black')
plt.show();
```

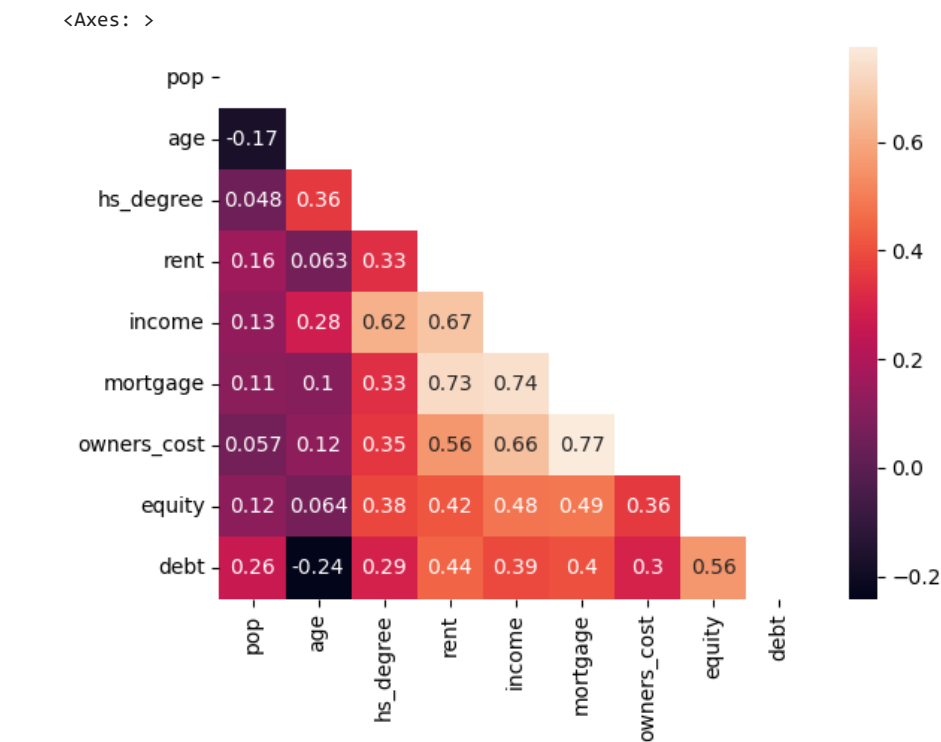


- The rent percentage are calculated diving the rent with family's income for the specific state.
- The overall USA rent percentage according to the family income of total population is 1.464%.
- The overall rent for Michigan state according the the family income of total population in Michigan is 1.387% which is lower than the overall country's rent percentage.
- The overall rent for New York according the the family income of total population in New York is 1.55% which is higher than the overall country's rent percentage.
- The overall rent for Indiana state according the the family income of total population in Indiana is 1.296% which is lower than the overall country's rent percentage.
- The overall rent for Wyoming state according the the family income of total population in Michigan is 1.094% which is lower than the overall country's rent percentage.

```
small_df= pd.DataFrame({
    'pop': df['pop'],
    'age': median_age_c,
    'hs_degree': df['hs_degree'],
    'rent': df['rent_median'],
    'income': df['family_median'],
    'mortgage': df['hc_mortgage_median'],
    # 'second_mortgage' : df['second_mortgage'],
    'owners_cost': df['hc_median'],
    'equity': df['home_equity'],
    'debt': df['debt']

})

corr= small_df.corr()
mask = np.triu(np.ones_like(corr, dtype=bool)) ## genrate a mask for upper triangle
sns.heatmap(small_df.corr(),cbar=True,annot=True,mask= mask)
```



From the above correlation heatmap we can observe that:

- There is a strong postive correlation between population and owners cost.
- There is negative correlation between individual's age and debt.
- There is strong postive correlation between highscool degree and individual's income.
- There is strong postive correlation between individual's income and rent.
- There is strong postive correlation between induvidual's rent and mortgage.
- There is strong postive correlation between individual's income and mortgage.
- There is strong postive correlation between mortgage and owners cost.
- There is postive correlation between equity and debt.

▼ Project Task 3:

Data Pre-processing:

```
pip install factor_analyzer

from factor_analyzer import FactorAnalyzer
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import FactorAnalysis
```



```
own = 1- df['debt']
```

```
lat_var= pd.DataFrame({
    'hs_degree' :df['hs_degree'],
    'median_age': median_age_c,
    'second_mortgage': df['second_mortgage'],
    'percent_own': own,
    'bad_debt': bad_debt})
```

```
id_data= ['hs_degree','median_age','second_mortgage','percent_own','bad_debt']
```

```
## standardizing the data
```

```
std= StandardScaler()
data_std= std.fit_transform(lat_var)
```

```
## Factor analysis
# Creating factor analysis object and perform factor analysis
```

```
n= len(id_data)
fa= FactorAnalyzer(n, rotation= 'varimax')
fa.fit(data_std)
```

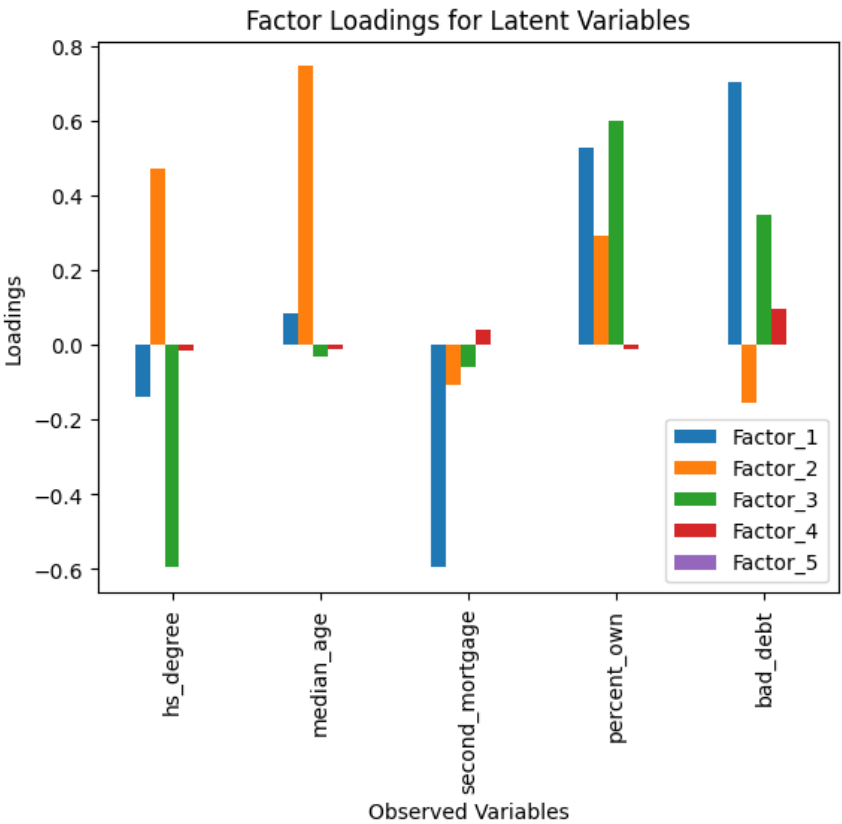
▼ FactorAnalyzer

FactorAnalyzer(n_factors=5, rotation='varimax', rotation_kwargs={})

```
loadings = pd.DataFrame(fa.loadings_, columns=[f'Factor_{i+1}' for i in range(n)], index=id_data)
```

```
## plotting the loadings of factor analysis
```

```
loadings.plot(kind='bar')
plt.title('Factor Loadings for Latent Variables')
plt.xlabel('Observed Variables')
plt.ylabel('Loadings')
plt.show()
```




```
# Check Eigenvalues
```

```
ev, v = fa.get_eigenvalues()
ev
```

```
array([2.07288437, 1.38923744, 0.72091405, 0.45958903, 0.35737511])
```

```
# Create scree plot using matplotlib
```

```
plt.scatter(range(1,lat_var.shape[1]+1),ev)
plt.plot(range(1,lat_var.shape[1]+1),ev)
plt.title('Scree Plot')
plt.xlabel('Factors')
plt.ylabel('Eigenvalue')
plt.grid()
plt.show()
```



Project Task 4:

Data modelling:

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score , mean_absolute_percentage_error, accuracy_score

## Selecting independent and dependent variable

y= df['hc_mortgage_mean']
x= df.drop(['hc_mortgage_mean', 'UID', 'COUNTYID', 'STATEID', 'state', 'state_ab', 'place', 'type', 'lat', 'lng', 'zip_code', 'area_code', 'city', 'place', 'primary'],axis=1)

## Split the data into training and testing sets

x_train, x_test, y_train, y_test= train_test_split(x,y)

## initialize Linear Regression Model

model= LinearRegression()

# Train the model on the training set
model.fit(x_train, y_train)

# LinearRegression
LinearRegression()

## making predictions on test set

y_pred= model.predict(x_test)

## Evaluate the model

r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)

print(f'R-squared: {r2:.2f}')
print(f'Mean Squared Error: {mse:.2f}')

R-squared: 0.99
Mean Squared Error: 4493.00

## Model Prediction at State level

def state_model(state):

## print("Model Prediction at State level \n\n")

state_df = df[df['state']== state]
x_state= state_df.drop(['hc_mortgage_mean', 'UID', 'COUNTYID', 'STATEID', 'state', 'state_ab', 'place', 'type', 'lat', 'lng', 'zip_code', 'area_code', 'city', 'place', 'primary'],axis=1)
y_state= state_df['hc_mortgage_mean']

X_train, X_test, Y_train ,Y_test= train_test_split(x_state, y_state)

model= LinearRegression()
model.fit(X_train, Y_train)
Y_pred= model.predict(X_test)

r2 = r2_score(Y_test, Y_pred)
mse = mean_squared_error(Y_test, Y_pred)
#acc= accuracy_score(Y_test, Y_pred)
mpe= mean_absolute_percentage_error(Y_test, Y_pred)

print(f'R-squared for {state}: {r2:.3f}')
print(f'Mean Squared error for {state}: {mse:.3f}')
#print(f'Accuracy score for {state}: {acc:.2f}')
print(f'Mean Absolute percentage error for {state}: {mpe:.4f}')
print('\n')

print("Model Prediction at State level \n\n")

for state in df['state'].unique() :
state_model(state)

Model Prediction at State level

R-squared for Michigan: 0.985
Mean Squared error for Michigan: 1897.614
Mean Absolute percentage error for Michigan: 0.0248

R-squared for Maine: 0.890
Mean Squared error for Maine: 11185.662
Mean Absolute percentage error for Maine: 0.0617

R-squared for Pennsylvania: 0.990
Mean Squared error for Pennsylvania: 2175.769
Mean Absolute percentage error for Pennsylvania: 0.0241

R-squared for Kentucky: 0.987
Mean Squared error for Kentucky: 1557.072
Mean Absolute percentage error for Kentucky: 0.0266

R-squared for Texas: 0.987
Mean Squared error for Texas: 4036.328
Mean Absolute percentage error for Texas: 0.0290

R-squared for Florida: 0.984
Mean Squared error for Florida: 4664.281
```

Mean Absolute percentage error for Florida: 0.0296

R-squared for Georgia: 0.984
Mean Squared error for Georgia: 3299.331
Mean Absolute percentage error for Georgia: 0.0305

R-squared for New York: 0.973
Mean Squared error for New York: 15890.515
Mean Absolute percentage error for New York: 0.0403


R-squared for California: 0.982
Mean Squared error for California: 6920.065
Mean Absolute percentage error for California: 0.0282

R-squared for Washington: 0.987
Mean Squared error for Washington: 3041.014
Mean Absolute percentage error for Washington: 0.0228

R-squared for Illinois: 0.990
Mean Squared error for Illinois: 3154.931
Mean Absolute percentage error for Illinois: 0.0242

```
df.to_csv('df.csv')
```

```
plt.hist(y_pred,bins=170)  
plt.show()
```



```
print('hello')
```

hello